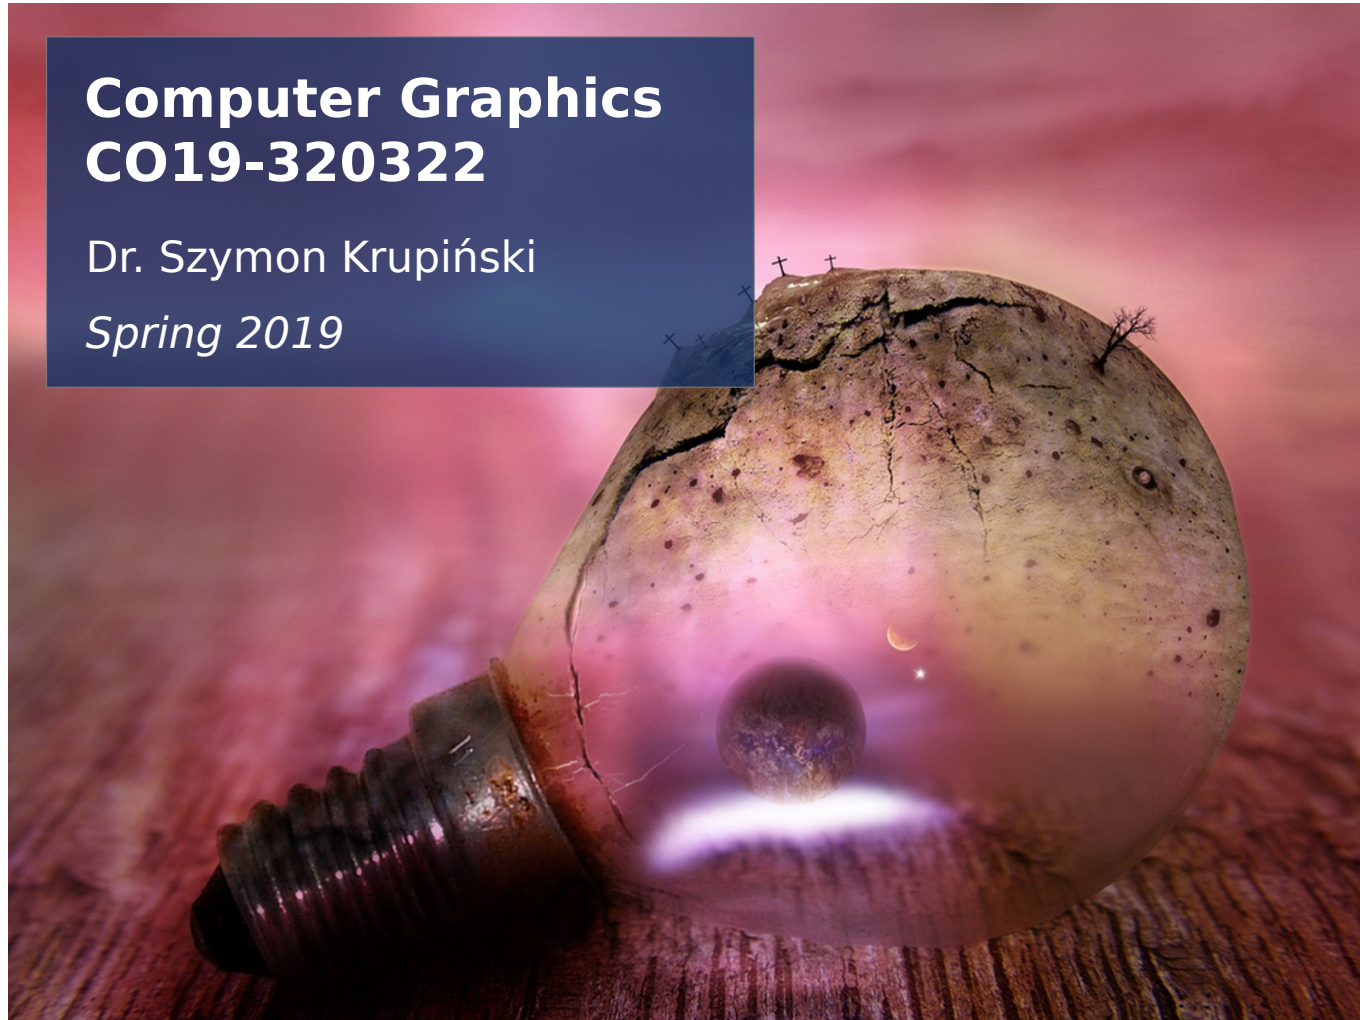


Lecture 4: Rasterisation 1


**Computer Graphics
CO19-320322**

Dr. Szymon Krupiński
Spring 2019




From yesterday: camera and projection

- A camera can be modelled using the pinhole model
- It has a number of important parameters which are often visualised as a camera frustum defining the geometry of the projection
- We constructed a mathematical expression for perspective transform which yields uniform coordinates. Good news, it's a matrix multiplication + homogeneous coordinates normalisation \rightarrow can be applied efficiently
- To be tested: what happens if we apply the projection to the frustum itself?




The unit box is where all the (visible) action happens!


Where are we now?




Modeling transforms:
Position object in scene



Viewing (camera) transform:
positions objects in coordinate
space relative to camera
Canonical form: camera at origin
looking down -z




**Projection transform +
homogeneous divide:**
Performs perspective
projection
Canonical form: visible
region of scene contained
within unit cube



Screen transform:
objects now in 2D screen coordinates


Screen transformation

- We didn't talk about the details of the last transformation because... it's rather simple! (but watch out for screen coordinate convention which differs from library to library):




Facing the triangles


- Up to now, we were transforming points, imagining wire meshes built upon them
- We had a first glance at a non-trivial problem: how to determine which pixels should be attributed to a given part of geometry
- Our concentration for today: project a triangle!



Facing the triangles



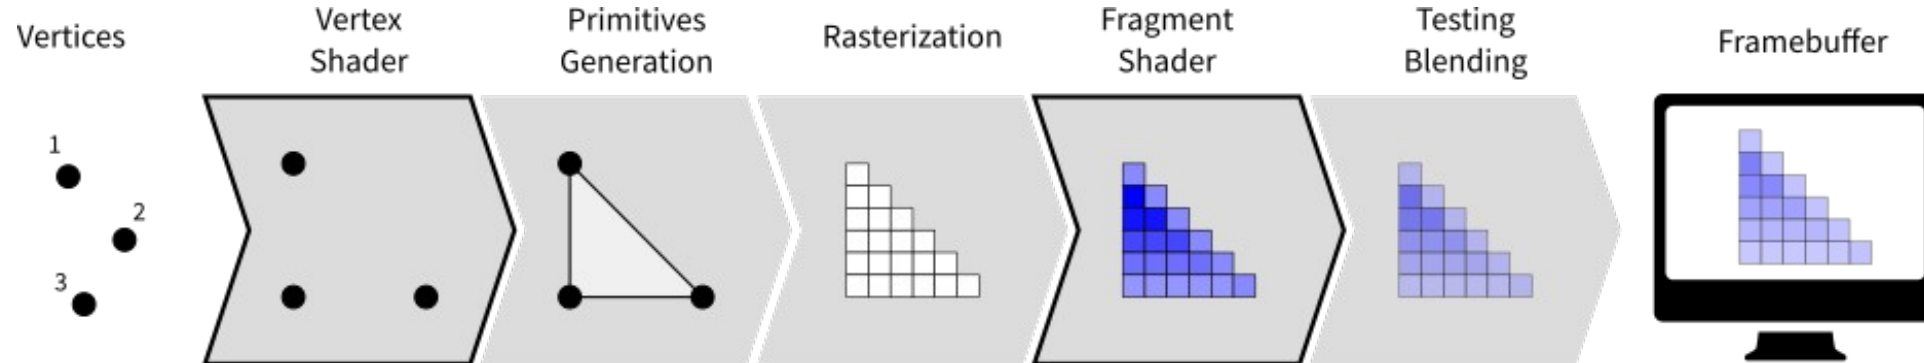
Question 3: how does one draw it realistically?



Question 4: how to cover all geometry and attribute all pixels on the screen?


Yet another broader look

- Without much ado, we are beginning to cover some steps of the modern CG pipeline, such as implemented in the recent versions of OpenGL



- We are currently looking at the stage called **rasterisation**, which is the task of taking an image described in a vector graphics format (shapes) – which we can obtain using projection - and converting it into a **raster image** (a series of pixels, dots or lines that when they come together on a display, they recreate the image).



Rasterisation or ray casting/tracing?



- Our transformation/projection chain allows us to methodically convert 3-D points to 2-D pixel locations
- But, in the class explanations, the metaphor of tracing light rays coming off/towards the objects is often in use; are we then using ray tracing?
- Rasterisation is considered an efficient and simple, albeit limited solution to **ray tracing**, which is supposed to mimic the behaviour of light and give high fidelity results




Coverage

- Which of the triangles actually “cover” the square pixel?
- Our objective is to colour the pixels according to the (arbitrary) colour of the triangle which they should represent
- We could potentially measure the percentage of the pixel occupied by every shape and colour it proportionally



Coverage


- Analytical schemas like this get much less convenient when we consider occlusions:



At this stage, it is interesting to reflect on a good formulation of this problem in terms of sampling theory!


Sampling – backbone of rasterisation

- Let's start with a 1-D signal which we want to convert to a discrete format (=rasterise).




Sampling – backbone of rasterisation

- Sampling is expensive, so we are trying to limit the number of samples



Sampling – backbone of rasterisation


- Trouble starts when we want to reconstruct our signal...



Sampling theorem is a fundamental bridge between continuous-time signals and discrete-time signals. It establishes a sufficient condition for a sample rate that permits a discrete sequence of samples to capture all the information from a continuous-time signal of finite bandwidth. The sampling theorem introduces the concept of a sample rate that is sufficient for perfect fidelity for the class of functions that are bandlimited to a given bandwidth.


Sampling – backbone of rasterisation

- Reconstruction: piecewise **constant** approximation



Sampling – backbone of rasterisation


- Reconstruction: piecewise **linear** approximation



None of them can overcome the natural limits due to Nyquist frequency!

Sampling – backbone of rasterisation

- Reconstruction: denser sampling can often improve the case



Sampling mathematics

- Dirac delta function $\delta(x)$ for $x \neq 0, \delta(x) = 0$ and $\int_{-\infty}^{\infty} \delta(x) dx = 1$
- We can use it to “pick” a value out of a continuous function $f(x)$:





$$\int_{-\infty}^{\infty} f(x) \delta(x) dx = f(0)$$

- We can construct a “comb” of shifted delta functions:

$$\text{III}_T(x) = T \sum_{i=-\infty}^{\infty} \delta(x - iT)$$

- ...and apply it to the function of interest:

$$\text{III}_T(x) f(x) = T \sum_{i=-\infty}^{\infty} f(iT) \delta(x - iT)$$




Sampling mathematics

- How do we handle reconstruction?
- Convolution


$$(f * g)(x) = \int_{-\infty}^{\infty} \underbrace{f(y)}_{\text{filter}} \underbrace{g(x-y)}_{\text{input signal}} dy$$

output signal
filter
input signal




- The filter can be quite simple, e.g. a box or a triangle function


$$f(x) = \begin{cases} 1 & |x| \leq 0.5 \\ 0 & \text{otherwise} \end{cases}$$



$$h(x) = \begin{cases} (1 - \frac{|x|}{T})/T & |x| \leq T \\ 0 & \text{otherwise} \end{cases}$$




- In practice: normalised sinc filter, or similar




Sampling: motivation

- We can treat the question of coverage (= which triangle covers which pixel) as a 2-D function: $\text{coverage}(x,y)$
 - Pixels define our sampling period
 - Here: the pixel center coordinate are used to sample the triangles coverage
- Plus, sampling theory will teach us a couple of lessons:
 - Subsampling helps reconstruct the signal
 - Aliasing (and anti-aliasing!)



Aliasing in the left
downsampled image –
can you tell which way
the lines are going?




Temporal
aliasing – the
wagon wheel
optical illusion



Occlusion

- Which triangle is in front of the other one?
 - ray tracing reply: whichever is hit by a ray sent out of the camera first
 - rasterisation reply: we can calculate it, if we want
 - it's an expensive operation, though
 - while doing screen projection for points in our shapes, we already have some z-depth information
 - It would be great to extend it to all points in the triangle – we would just need to interpolate between the vertices!
 - We can construct an image with this information for all pixels and keep track when new elements come
→ **z-buffer**



This is NOT an innocent remark! We will spend some time looking at this problem!

Thank you!

- Questions?

